

Peer to Peer Networks

Jan H. Knudsen (20092926) Roland L. Pedersen (20092817)
Kris V. Ebbesen (20094539)

May 18, 2014

Contents

1	Introduction	3
2	Related Work	3
3	Use Case	3
3.1	Use case 1: Chinese Dissidents	3
3.2	Use case 2: Young Love	4
4	A short note on cryptography	5
5	Method of Operation	5
5.1	Base System	5
5.2	Encrypting Messages and Hiding Recipients	5
5.3	Signed Messages and Acknowledgement	6
5.4	Encrypting Peer Communication	6
5.5	Cover Traffic	7
5.6	Providing Proof of Work	7
5.7	Public Key Distribution	7
5.8	Limiting RPC Call Availability	8
6	Manual of Operations	8

1 Introduction

With the recent rise in awareness about the blanket surveillance of our daily internet usage, the general population no longer feel that the anonymity and security of their online exchanges are guaranteed. We hope to amend this.

In this project report, we describe how to develop a peer to peer system that allows the exchange of encrypted chat messages, while making it quite difficult to determine who sent and received what messages. The system is built on top of an existing unstructured network, and the techniques we use apply to many different unstructured networks, this means, that with minor extensions, existing networks can be allowed to support this kind of traffic. We will describe both why we have chosen to develop this system, how it works, how to operate it, and show experiments that measure the performance of the network.

2 Related Work

Protecting freedom of speech and the free exchange of ideas in the face of government oppression and censorship is one of the nobler goals of p2p networks and cryptography. Several systems and protocols with this as their main goal have been designed and analysed over the years. One such project is Freenet (Clarke et. al, 2002) which builds a virtual file space distributed among peers and provides censorship resistance by being anonymous, decentralized, encrypted with peers only knowing about direct neighbours, and providing peers plausible deniability by having the shared resources be encrypted so forwarding peers can't know the content. Freenet, however, provides very low performance and no guarantee that a file will remain, as replication relies on the resource being requested by other peers and other peers choosing to cache it along the way. Resources are also available to anyone, meaning encryption of shared resources and another channel to distribute the keys is paramount. Resources also cannot be sent to a specific target.

A different system that provides something closer to a realtime instant message client is BitMessage(Warren, 2012). This, however, uses flooding to send messages across the network and as we have demonstrated in our previous paper (Ebbesen, Knudsen and Pedersen, 2014), there are better performing alternatives, such as random k-walkers, and we would like to see if we could implement something as secure as BitMessage using these different methods of sending messages.

3 Use Case

When developing the system we maintained clear goals for the required features of the system, which helped us not only remember why we were developing the system, but also guide the project in the right direction.

These goals are best illustrated in use case format, as follows:

3.1 Use case 1: Chinese Dissidents

Primary Actor: 2 Chinese citizens with a wish for democracy.

Goal: To exchange thoughts and literature about democracy in a way that is safe, secure, and not under government scrutiny.

Other interested Parties: The Chinese government.

Preconditions: It is assumed that both citizens have exchanged public keys prior to the use case (by carrier pigeon). In order to avoid data analysis attempts, a secure channel to a peer located outside the reach of the Chinese government would be also preferred.

Success Criteria: Using the developed system, the two citizens should be able to send chat messages to each other, with a guarantee that they cannot be read by 3rd party actors. Additionally, they should be able to know who sent the messages, and whether their messages reach their destination. Finally, it should be difficult to determine that these two citizens are communicating.

Method:

1. Both citizens add their private key to the system.
2. Both citizens set up the public key of each other in the system.
3. The citizens start sending messages.
4. The citizens look for confirmation messages from the system.

3.2 Use case 2: Young Love

Primary Actor: The young doe-eyed girl Earlene.

Goal: Earlene wishes to profess her love to Billy Ray, the handsome new farmhand, in intimate prose, yet she does not wish to reveal her identity yet.

Other interested Parties: Billy Ray, the manliest farmhand around.

Preconditions: It is assumed that Earlene has obtained the hash address of Billy Ray, either through teenage girl gossip, or to phone book. Also required is Billy Ray's willing publishing of his public key.

Success Criteria: Earlene should be able to express her innermost desires to Billy Ray, without him know her identity.

Method:

1. Earlene uses the system to fetch the public-key of Billy Ray.
2. Earlene condenses her desires into a 20-page novella with graphic descriptions of the surrounding landscape.
3. Earlene sends the message unsigned to Billy Ray.
4. The message is received at Billy Rays end, and Earlene receives a signed acknowledgement.

Additional details: Should Earlene wish to continue secret communication with Billy Ray, she might generate a new RSA key pair, publish the public key, and attach the corresponding hash address in her message.

4 A short note on cryptography

The project described in this report relies heavily on cryptographic system and practices in order to be possible. It is however the case that none of the persons working on the project have any previous experiences working with secure systems, since we come from backgrounds in algorithms and computer graphics. We have chosen such a heavy reliance on systems outside our normal line of work as a learning experience, and with the strong conviction that most cryptographic systems work well as black boxes.

Expanding further on this, when we refer to a cryptographic method, we will rarely expand on the inner workings of such components, but rather rely on their security as provided by their developers. Of course this makes us somewhat prone to making errors that would be considered mistakes by hardened security veterans, but we beg forgiveness for our bright-eyed naïveté.

In the end, what matters to us in this project, is the parts that relate to peer-to-peer systems.

5 Method of Operation

5.1 Base System

The system described is built atop the unstructured network developed during the P2PN course (TODO: ref earlier P2PN report). This network contains very little structural information, and bases its topology on the GIA network (Chawathe et al., 2003).

The choice of this network was made based on its simplicity and extendibility, and due to the fact that unstructured networks require little information about the peers involved, making it difficult to track which peers are doing what.

Note that the techniques used to extend the network could be applied to most unstructured networks, and would probably work just as well on the GIA network.

5.2 Encrypting Messages and Hiding Recipients

In order to ensure that no adversaries can read the content of any given message, we encrypt chat messages travelling across the network using RSA-OAEP (Bellare and Rogaway, 1995). RSA-OAEP was chosen due to its ease of use, and security against repeated plaintext attacks.

When performing this encryption, we use a pair of RSA (TODO:Mayby reference?) keys. The sender must obtain the public key of the final recipient (how to do this will be explained later), in order to encrypt the message.

When the chat message is sent, it is first encrypted by RSA-OAEP using the public key of the recipient, and then broadcast across the network using either flooding or k-walkers. Whenever a peer receives a messages travelling across the network, it will attempt to decrypt it using the corresponding RSA-OAEP decryption using its own private key. This will fail for all peers except the recipient, ensuring that only the final recipient will be able to obtain the contents of the chat message.

Note that the encrypted message sent across the network contains no delivery address of any kind, and as such no other peers will know the final recipient.

It is also worth noting that only one RSA key pair is required to send messages. The sender needs no private key, nor do any other peers in the network except the receiver.

5.3 Signed Messages and Acknowledgement

All chat messages in the system may or may not be signed by the sender. If the sender wishes not to sign his messages, in order to hide his identity from the receiver, or because he is not in possession of a private key, he may omit this signature. Additionally, any message received by a peer can be acknowledged by returning a signed digest of the message.

Both types of signatures are done according to **PKCS#1 v1.5** (TODO:Some sort of reference).

In the case of the sender signing a message, we send a signature of the plain-text message along with the encrypted message. This ensures, that only after obtaining the decrypted message will it be possible to verify the signature. This ensures that we keep the identity of the sender hidden to anyone except the recipient, and that the recipient can securely verify the sender given his public key. Also, should anyone attempt to tamper with the message before delivery, the signature will no longer be valid.

When verifying the delivery of a message the receiver returns a signed digest of the plaintext message, which is verified by the sender. This ensures that the sender has received the message, as he is the only one able to provide a valid signature. If the peer is using flooding we simply return this value as part of the xml-rpc call, while we answer back using a k-walker in the case that we receive a message by k-walker. Given a small random delay, it becomes difficult to determine whether a message was received by any given peer, or one of his neighbours. Also, should anyone tamper with the message before delivery, the receiver will no longer sign the correct data, making this easily detectable for the sender.

5.4 Encrypting Peer Communication

All traffic between peers in the network is encrypted using anonymous Diffie-Hellman (Diffie and Hellman, 1976) encryption. This encryption is provided by wrapping connections between peers in an SSL layer, with no certificates and anonymous Diffie-Hellman as the only cipher set.

This ensures that peers can communicate without outside parties snooping on the information, which makes it very hard to track messages across the network, since the data sent from messages, cover traffic, and general networks operations will be indistinguishable.

Another reason to use anonymous Diffie-Hellman encryption is that it enforces no requirements on previously distributed keys or identities of the peers, keeping each peer's knowledge about its neighbours at a minimum.

In order to prevent constant Diffie-Hellman key renegotiations we provide cached pools of SSL connections, meaning that we only create a new connection when the peer runs out of idle connections to the same peer.

5.5 Cover Traffic

Inspired by the use of cover traffic in the Tarzan p2p protocol (Freedman and Morris, 2002), we include cover traffic in our solution to preventing traffic analysis. The network relies heavily on the SSL encryption of the peer-to-peer connections to keep traffic types indistinguishable, making it very difficult for an outside observer to discern what data traffic belongs to messages and which concern the network.

In terms of cover traffic, we provide two sources of cover.

One is the general operations of the network. Neighbours will constantly contact each other to ensure that they are alive, and any peers leaving or joining the network will require a fair bit of communication between peers. Since all of this traffic is encrypted, it will be hard to distinguish this communication from messages.

The second source of cover traffic is explicit cover traffic. Peers will at random intervals send random data of random lengths between each other. This data ensures unpredictable network traffic, and hinders traffic analysis even further.

5.6 Providing Proof of Work

To keep the network stable, and free from Sybil-style attacks, we use a proof of work system. This ensures that peers that wish to put a strain on the network, or affect the overlay network structure, will need to expend large amounts of computational resources to do so.

The proof of work system is based on HashCash (Back, 2002), and requires a peer to generate a partial hash collision with the timestamped resource, using the SHA-256 hashing algorithm. How large a collision and how new a time stamp is fully configurable.

A proof of work is currently required in 2 circumstances.

The first is when a peer wishes to join the neighbourhood of another peer. In requiring a proof of work for joining or moving within the network, we make Sybil and Eclipse attacks less likely, while imposing little to no hindrance on long-term stable peers.

The second proof of work is required when a peer wishes to send a message. This is to deter spamming of the network, and to prevent malicious peers from forcing other peers to spend an unwanted amount of time trying to decrypt messages, or drown a single peer in messages after having obtained its public key.

Note that the standard settings for the required bits of a proof of work are currently quite low, in order to allow rapid testing of the network

5.7 Public Key Distribution

When a peer wishes to communicate chat messages to another peer, it is required to know the public key of the recipient.

This public key can be supplied directly by the sender, indicating that the key has been distributed securely outside of the network. In this case, the key is simply loaded from a provided file.

The network also offers the option of publishing public keys using the underlying peer-to-peer network's ability to share resources. When doing this, the public key is read, and stored in the network as a resource using the base64 encoding of its SHA-256 hash

as its name. The key can then either be fetched and stored normally as a resource by other peers, or loaded directly into the public storage of other peers.

Any peer that loads the key directly will verify its hash as it does so.

The result is a tag (44 characters long), that can be shared much easier than an entire public key.

5.8 Limiting RPC Call Availability

Standard practice in object-oriented Python-based RPC servers is to register the entire object for RPC call availability. This is highly inadvisable if one wishes to protect the network from malicious peers.

In order to prevent this form of attack, we enforce strict limitations of function availability. This is done by extending the way the RPC calls are handled by the XML-RPC components, and tagging only the needed methods calls as being callable by RPC. Any attempt to call an unlisted function will silently be ignored.

6 Manual of Operations

In order to operate a peer in the network, one must rely on either python scripting against the peer class of the source code, or the supplied command line interface.

We here explain the commands required to operate the peer using the command line interface. Scripting directly against peer class is left as an exercise for the reader.

Note that the peer still supports most of the commands of the original network (TODO:Ref P2PN paper).

The commands are as follows:

hello [*address*] Attempt to join the network. An optional address parameter may be specified in order to bootstrap against a known peer.

After joining the network, the peer will be ready to add keys, and chat. Please note that it might take several seconds for the peer to establish an acceptable amount of neighbours.

secret *private_key* Load the given private key from a local file, and set is as the current key used for decrypting and signing messages. This is required in order to receive messages encrypted with the corresponding public key, and to sign messages sent from the local peer. Note that the key must be an RSA private key in the *pem* format.

friend *name* *public_key* Load the given public key from a local file, and associate it with the alias provided by the name parameter. This is required in order to send messages to the peer with the corresponding private key, and to identify that peer as a sender. Note that the key must be an RSA public key in the *pem* format.

publish *public_key* Make the given public available for retrieval through the peer to peer network. Shortly after entering this command the peer will display a hash of your key, which you can share. This allows other peers to download your public

key through the network if they have the corresponding hash string. Note that the key must be an RSA public key in the *pem* format.

friend *name hash* Fetches the key stored in the network under the given hash, and checks for hash validity of the key. If successful, the public key retrieved will be associated with the alias specified in the name parameter.

message *name message* Attempts to deliver a message to a friend added under the alias specified by the name parameter, with the content of the message parameter, using flooding. The message will be signed if possible, and a report of delivery given.

kmessage *name message* Attempts to deliver a message to a friend added under the alias specified by the name parameter, with the content of the message parameter, using k-walkers. The message will be signed if possible, and an id will be given, which allows matching to a later received acknowledgement.

References

- Yatin Chawathe, Sylvia Ratnasamy, Lee Breslau, Nick Lanham, and Scott Shenker. Making gnutella-like p2p systems scalable. In *Proceedings of the 2003 Conference on Applications, Technologies, Architectures, and Protocols for Computer Communications*, SIGCOMM '03, pages 407–418, New York, NY, USA, 2003. ACM. ISBN 1-58113-735-4. doi: 10.1145/863955.864000. URL <http://doi.acm.org/10.1145/863955.864000>.
- Mihir Bellare and Phillip Rogaway. Optimal asymmetric encryption. In Alfredo Santis, editor, *Advances in Cryptology - EUROCRYPT 94*, volume 950 of *Lecture Notes in Computer Science*, pages 92–111. Springer Berlin Heidelberg, 1995. ISBN 978-3-540-60176-0. doi: 10.1007/BFb0053428. URL <http://dx.doi.org/10.1007/BFb0053428>.
- Whitfield Diffie and Martin E Hellman. New directions in cryptography. *Information Theory, IEEE Transactions on*, 22(6):644–654, 1976.
- Adam Back. Hashcash - A Denial of Service Counter-Measure. Tech Report, 2002 URL <http://www.hashcash.org/papers/hashcash.pdf>
- Freedman, Michael J. and Morris, Robert. Tarzan: a peer-to-peer anonymizing network layer, 2002
- Clarke, Ian and Miller, Scott G. and Hong, Theodore W. and Sandberg, Oskar and Wiley, Brandon. Protecting Free Expression Online with Freenet, 2002
- Jonathan Warren Bitmessage: A Peer-to-Peer Message Authentication and Delivery System URL <https://bitmessage.org/bitmessage.pdf>
- Kris V. Ebbesen, Jan H. Knudsen and Roland L. Pedersen. Peer to Peer Networks. Course Report, 2014